

Project Name : Large Scale Wave Energy Farm :

Submitted by : Veenalakshmi P V

Date: 28-10-2024

**Table of Contents : **

1. Introduction
2. Goal Of the Project
3. Data Story
4. Data Preprocessing
5. Data Cleaning
6. Data transformation
7. Handling Outliers
8. Data Splitting
9. Data Visualisation
10. Feature Importance
11. Model Selection
12. Feature Selection
13. Pipeline
14. Model Evaluation and Predictions
15. Hyper parameter tuning
16. Saving the Model
17. Load the Model

Introduction :

This project has significance in renewable energy research, as accurately predicting power output can help in optimizing the design and operation of wave energy converters, leading to more efficient and sustainable energy production.

****Goal Of this Project : ****

The primary goal of this project is to analyze and compare power consumption patterns between Perth and Sydney across two datasets with varying regional granularity: Perth_49 vs. Sydney_49 and Perth_100 vs. Sydney_100. Using machine learning regression models (e.g., Random Forest, Linear Regression), the project aims to:

1.Predict Power Consumption:

Create predictive models to accurately estimate power usage based on features provided in the dataset.

Identify Key Factors: Use feature importance analysis to understand which factors most influence power consumption in each region.

2.Detect Patterns and Differences:

there are significant differences in power usage patterns between Perth and Sydney and assess whether these vary with the dataset's regional granularity (49 vs. 100 regions).

3.Evaluate Model Performance and Overfitting:

Test models on both training and testing sets to ensure generalizability and prevent overfitting, allowing the model to perform well on unseen data.

4.Provide Insights for Energy Management:

Derive insights that could potentially help energy providers or policymakers in planning and optimizing power distribution in Perth and Sydney.

2. Data Story :

Comparing Power Consumption in Perth and Sydney :

In this data story, Analyzing the power consumption data from two Australian cities: Perth and Sydney. The datasets, labeled Perth_49, Sydney_49, Perth_100, and Sydney_100, represent energy consumption measurements from these cities over time. Each dataset contains information on power usage across different time intervals or features. Our goal is to extract insights into power consumption patterns, detect anomalies or outliers, and highlight key features that contribute to variations in energy usage.

Objective:

The objective of this analysis is to:

Compare power consumption between Perth and Sydney.

Detect outliers in the data using the Z-score method.

Visualize patterns in power usage using time-series and distribution plots.

Identify key features contributing to power consumption through feature importance analysis.

Data Preprocessing Step:

Load The Data Set :

```
In [5]: import pandas as pd
import numpy as np

perth_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
perth_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
sydney_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WE
sydney_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\W

print(f"\nPerth_49 : {perth_49.head()}, \n\nPerth_100 : {perth_100.head()}, \n\
```

Perth_49 :		X1	Y1	X2	Y2	X3	Y3	X4	Y4	X5
Y5 ... \										
0	600.0	0.0	546.16	37.50	489.79	74.88	432.47	112.05	650.0	0.0 ...
1	593.0	12.0	546.16	37.50	489.79	74.88	432.47	112.05	644.0	8.0 ...
2	593.0	12.0	546.16	37.50	489.79	74.88	432.47	112.05	644.0	8.0 ...
3	593.0	12.0	546.16	37.50	489.79	74.88	432.47	112.05	644.0	8.0 ...
4	200.0	0.0	146.17	37.53	89.76	74.93	32.40	112.18	400.0	0.0 ...

	Power42	Power43	Power44	Power45	Power46	Power47	Power48	\
0	88867.92	98844.30	101283.59	98934.63	101624.58	100915.03	99625.68	
1	88896.55	98759.79	101346.07	98873.59	101629.01	100934.53	99606.13	
2	88919.83	98746.68	101346.15	98875.57	101618.32	100941.00	99611.35	
3	88855.14	98760.96	101338.59	98971.58	101632.28	100943.59	99589.25	
4	88005.30	98630.24	100432.73	98803.01	101064.48	100948.38	99028.87	

	Power49	qW	Total_Power
0	96704.34	0.87	4102461.43
1	96718.39	0.87	4103361.41
2	96719.14	0.87	4103680.44
3	96735.04	0.87	4105661.06
4	96286.71	0.79	3752648.77

[5 rows x 149 columns],

Perth_100 :		X1	Y1	X2	Y2	X3	Y3	X4	Y4
X5 Y5 \									
0	1000.0	0.0	946.08	37.42	889.67	74.76	832.02	112.10	1250.0 1.0
1	800.0	0.0	746.04	37.38	689.81	74.79	632.43	111.97	1200.0 0.0
2	600.0	0.0	545.98	37.50	489.87	74.95	432.52	112.15	650.0 0.0
3	600.0	0.0	546.09	37.41	489.70	74.76	432.35	111.90	800.0 0.0
4	600.0	0.0	545.95	37.52	489.69	74.99	432.46	112.16	800.0 0.0

	...	Power93	Power94	Power95	Power96	Power97	Power98	\
0	...	98711.68	102872.82	100743.44	99259.87	98909.46	101388.37	
1	...	96351.38	102253.02	101744.20	99482.45	99304.59	101953.23	
2	...	96985.29	102472.80	101757.63	99429.45	98709.14	101312.44	
3	...	76823.85	88005.41	98779.92	100260.30	98744.25	101144.58	
4	...	98195.95	102874.16	100256.15	99221.12	98969.54	101389.77	

	Power99	Power100	qW	Total_Power
0	101025.35	98676.66	0.75	7257985.04
1	100878.42	99508.49	0.74	7103374.61
2	100979.86	99024.16	0.76	7335380.64
3	100835.43	98915.38	0.75	7187769.87
4	100924.02	98796.00	0.75	7260222.61

[5 rows x 302 columns],

Sydney_49 :		X1	Y1	X2	Y2	X3	Y3	X4	Y4	X5
Y5 ... \										
0	1.0	1.0	0.00	70.00	1.00	140.0	50.00	198.0	401.0	1.0 ...
1	598.0	0.0	595.82	77.69	593.70	150.0	549.00	198.0	798.0	0.0 ...
2	198.0	0.0	197.46	75.19	192.94	150.0	87.64	198.0	398.0	0.0 ...
3	598.0	0.0	596.97	69.41	592.69	143.8	549.00	198.0	398.0	200.0 ...
4	198.0	0.0	197.18	79.83	192.97	150.0	89.53	198.0	398.0	0.0 ...

	Power42	Power43	Power44	Power45	Power46	Power47	Power48	\
0	71909.82	70674.49	70972.33	90957.03	90903.63	87876.82	79499.23	
1	68757.68	70665.50	69963.48	84511.25	85691.70	85211.51	76678.20	
2	73675.80	77808.44	73519.82	91436.35	88770.60	86632.78	77932.46	
3	68947.21	71668.05	69380.67	85191.27	84453.12	85300.41	78573.25	

```
4 78367.97 79075.06 74354.03 85254.75 86978.69 86951.65 77671.87
```

```
      Power49      qW  Total_Power
0 68880.39 0.78 4065416.61
1 76119.53 0.76 3951216.37
2 69343.12 0.78 4022640.78
3 72527.16 0.75 3879223.41
4 74901.38 0.77 3974691.24
```

```
[5 rows x 149 columns],
```

```
Sydney_100      X1  Y1      X2      Y2      X3      Y3      X4      Y4      X5
Y5 ... \
0 1.0 1.0 1.00 51.00 1.00 101.00 1.00 151.0 398.0 0.0 ...
1 198.0 0.0 197.18 80.53 193.59 150.00 77.58 198.0 598.0 0.0 ...
2 198.0 0.0 197.07 76.64 192.74 155.74 84.67 198.0 798.0 0.0 ...
3 1.0 1.0 1.00 51.00 1.00 101.00 1.00 151.0 398.0 0.0 ...
4 198.0 0.0 197.46 75.07 197.18 149.14 149.00 198.0 598.0 0.0 ...
```

```
      Power93      Power94      Power95      Power96      Power97      Power98      Power99 \
0 74018.52 71727.79 67966.45 63101.26 88826.02 86531.44 83786.68
1 63702.46 67776.99 65133.52 63138.74 82852.91 83519.30 81973.65
2 55788.34 59593.98 60073.60 59198.12 63377.08 72078.85 77435.62
3 66961.48 65716.93 66637.89 62562.54 80858.08 82656.53 82171.28
4 51814.27 59556.86 68341.92 70731.90 64192.86 69757.10 75581.40
```

```
      Power100      qW  Total_Power
0 73514.19 0.69 7247491.41
1 71781.34 0.67 7119352.90
2 67457.26 0.68 7148342.69
3 71713.30 0.69 7317998.83
4 69741.63 0.65 6925096.49
```

```
[5 rows x 302 columns]
```

```
# info()
```

```
In [7]: sydney_100_data = pd.read_csv('C:\\Users\\Admin\\Desktop\\Large Scale Wave Ener
print("Sydney 100 Dataset Info:")
print(sydney_100_data.info())

# For Perth 49 dataset
perth_49_data = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy
print("\\n\\nPerth 49 Dataset Info:")
print(perth_49_data.info())

sydney_49_data = pd.read_csv('C:\\Users\\Admin\\Desktop\\Large Scale Wave Ener
print("\\n\\nSydney 49 Dataset Info:")
print(sydney_49_data.info())

perth_100_data = pd.read_csv('C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy
print("\\n\\nPerth 100 Dataset Info:")
print(perth_100_data.info())
```

```
Sydney 100 Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2318 entries, 0 to 2317
Columns: 302 entries, X1 to Total_Power
dtypes: float64(302)
memory usage: 5.3 MB
None
```

```
Perth 49 Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36043 entries, 0 to 36042
Columns: 149 entries, X1 to Total_Power
dtypes: float64(149)
memory usage: 41.0 MB
None
```

```
Sydney 49 Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17964 entries, 0 to 17963
Columns: 149 entries, X1 to Total_Power
dtypes: float64(149)
memory usage: 20.4 MB
None
```

```
Perth 100 Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7277 entries, 0 to 7276
Columns: 302 entries, X1 to Total_Power
dtypes: float64(302)
memory usage: 16.8 MB
None
```

Describe ()

```
In [8]: import pandas as pd
import numpy as np

perth_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
perth_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
sydney_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WE
sydney_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\W

# For Sydney 100 dataset
print("\n\nSydney 100 Dataset Summary:")
print(sydney_100_data.describe())

print("\n\nSydney 49 Dataset Summery:")
print(sydney_49_data.describe())

print("\n\nPerth 100 Dataset Summery:")
print(perth_100_data.describe())

print("\n\nPerth 49 Dataset Summery:")
print(perth_49_data.describe())
```

25%	50.000000	500.000000	50.000000	800.000000	0.000000
50%	74.820000	500.000000	100.000000	700.000000	0.080000
75%	74.960000	632.750000	112.150000	850.000000	50.000000
max	990.000000	1000.000000	990.000000	1000.000000	919.590000

	...	Power42	Power43	Power44	Power45	\
count	...	36043.000000	36043.000000	36043.000000	36043.000000	
mean	...	93678.772248	96530.68484	96666.293181	97007.214249	
std	...	7401.226140	6709.53446	7020.690028	4829.877255	
min	...	52516.130000	56391.97000	53877.360000	53050.330000	
25%	...	88177.210000	94648.08000	96932.520000	97612.350000	
50%	...	93694.540000	98729.91000	99269.310000	98857.150000	
75%	...	100997.520000	100622.52000	100282.360000	99156.130000	
max	...	110945.940000	109400.43000	114194.520000	106702.150000	

		Power46	Power47	Power48	Power49	\
count	36043.000000	36043.000000	36043.000000	36043.000000		
mean	98466.265281	98106.278501	97462.663041	96134.920454		
std	4978.194259	4263.508074	3134.420742	3889.098339		
min	55401.380000	63028.260000	61717.310000	47257.430000		

Handle Missing Data


```
In [9]: # Checking for missing values in all datasets
print("Missing values in WEC_Sydney_49:\n", sydney_49_data.isnull().sum()) # Ch
print("Missing values in WEC_Sydney_100:\n", sydney_100_data.isnull().sum()) #
print("Missing values in WEC_Perth_49:\n", perth_49_data.isnull().sum()) # Chan
print("Missing values in WEC_Perth_100:\n", perth_100_data.isnull().sum()) # Ch
```

```
Missing values in WEC_Sydney_49:
  X1      0
  Y1      0
  X2      0
  Y2      0
  X3      0
  ..
Power47   0
Power48   0
Power49   0
qW        0
Total_Power  0
Length: 149, dtype: int64
Missing values in WEC_Sydney_100:
  X1      0
  Y1      0
  X2      0
  Y2      0
  X3      0
  ..
Power98   0
Power99   0
Power100  0
qW        0
Total_Power  0
Length: 302, dtype: int64
Missing values in WEC_Perth_49:
  X1      0
  Y1      0
  X2      0
  Y2      0
  X3      0
  ..
Power47   0
Power48   0
Power49   0
qW        0
Total_Power  0
Length: 149, dtype: int64
Missing values in WEC_Perth_100:
  X1      0
  Y1      0
  X2      0
  Y2      0
  X3      0
  ..
Power98   0
Power99   0
Power100  0
qW        0
Total_Power  0
Length: 302, dtype: int64
```

```
In [ ]: This data set have no missing values
```

```
In [ ]: The output shows the results of checking each dataset for missing values. Here'

WEC_Sydney_49 Dataset: :
There are 149 columns in total. No columns contain missing values, as each colu

WEC_Sydney_100 Dataset:
This dataset has 302 columns, which is consistent with the inclusion of additio

WEC_Perth_49 Dataset:
Like WEC_Sydney_49, this dataset has 149 columns. No missing values are present

WEC_Perth_100 Dataset:
This dataset has 302 columns, with Power1 to Power100 included. All columns hav

Key Takeaways No Missing Data: Across all four datasets (WEC_Sydney_49, WEC_Syd

Data Completeness: The absence of missing values ensures that can directly proc
```

Handle Duplicates :

- o Check for duplicate rows using duplicated().
- o Remove duplicates: drop_duplicates()

```
In [10]: # Check for duplicate rows in Perth and Sydney datasets
perth_49_duplicates = perth_49.duplicated()
sydney_49_duplicates = sydney_49.duplicated()
perth_100_duplicates = perth_100.duplicated()
sydney_100_duplicates = sydney_100.duplicated()

# Print the number of duplicates found
print(f"\nPerth 49 Duplicates: {perth_49_duplicates.sum()}")
print(f"\nSydney 49 Duplicates: {sydney_49_duplicates.sum()}")
print(f'\nPerth 100 Duplicates:{perth_100_duplicates.sum()}')
print(f'\nSydney 100 Duplicates:{sydney_100_duplicates.sum()}')
```

Perth 49 Duplicates: 25107

Sydney 49 Duplicates: 13148

Perth 100 Duplicates:4540

Sydney 100 Duplicates:1084

In [4]: `import pandas as pd`

```
perth_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
perth_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
sydney_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WE
sydney_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\W

# Remove duplicates from both Perth and Sydney datasets
perth_49_cleaned = perth_49.drop_duplicates()
sydney_49_cleaned = sydney_49.drop_duplicates()
perth_100_cleaned = perth_100.drop_duplicates()
sydney_100_cleaned = sydney_100.drop_duplicates()

# Check the shape to confirm the removal
print(f"Perth 49 Shape after removing duplicates: {perth_49_cleaned.shape}")
print(f"Sydney 49 Shape after removing duplicates: {sydney_49_cleaned.shape}")
print(f"Perth 100 shape after removing duplicates: {perth_100_cleaned.shape}")
print(f"sydney 100 shape after removing duplicates: {sydney_100_cleaned.shape}")
```

Perth 49 Shape after removing duplicates: (10936, 149)
Sydney 49 Shape after removing duplicates: (4816, 149)
Perth 100 shape after removing duplicates: (2737, 302)
sydney 100 shape after removing duplicates: (1234, 302)

Data Transformation**

- o Encode categorical variables: Use `pd.get_dummies()` or `LabelEncoder`.
- o Scale numerical features: Apply `StandardScaler` or `MinMaxScaler`.

Data transformation is an essential part of this project, aiming to prepare raw datasets for model training, analysis, and visualization. Here's a structured approach to the transformations applied to the Perth and Sydney datasets in this power consumption comparison project

In [6]: `import pandas as pd`

```
# Load dataset from a CSV file into a DataFrame
perth_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
perth_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
sydney_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WE
sydney_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\W

perth_49_encoded = pd.get_dummies(perth_49, drop_first=True) # drop_first=True
perth_100_encoded = pd.get_dummies(perth_100, drop_first=True)
sydney_100_encoded = pd.get_dummies(sydney_100, drop_first=True)
sydney_49_encoded = pd.get_dummies(sydney_49, drop_first=True)
```

In []:

Scale Numerical Features :

Scaling ensures that all numerical features are on the same scale, improving model performance.

Using StandardScaler (standardization):

this code handles the scaling of numerical features in the specified datasets to prepare them for analysis or modeling. The choice of scaling method depends on the specific requirements of your analysis, such as whether to bound the features to a specific range or normalize them based on their distribution.

```
In [20]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler, StandardScaler

# Load the datasets
perth_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
perth_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
sydney_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WE
sydney_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\W

# Initialize scalers
min_max_scaler = MinMaxScaler()
standard_scaler = StandardScaler()

# Scale features for Sydney_49 vs Perth_49
sydney_49_scaled = min_max_scaler.fit_transform(sydney_49.select_dtypes(include=
perth_49_scaled = min_max_scaler.fit_transform(perth_49.select_dtypes(include=

# Scale features for Sydney_100 vs Perth_100
sydney_100_scaled = standard_scaler.fit_transform(sydney_100.select_dtypes(incl
perth_100_scaled = standard_scaler.fit_transform(perth_100.select_dtypes(includ

# Convert scaled data back to DataFrame
sydney_49_scaled_df = pd.DataFrame(sydney_49_scaled, columns=sydney_49.columns)
perth_49_scaled_df = pd.DataFrame(perth_49_scaled, columns=perth_49.columns)
sydney_100_scaled_df = pd.DataFrame(sydney_100_scaled, columns=sydney_100.colum
perth_100_scaled_df = pd.DataFrame(perth_100_scaled, columns=perth_100.columns)

# Optional: Save the scaled dataframes to CSV
sydney_49_scaled_df.to_csv('sydney_49_scaled.csv', index=False)
perth_49_scaled_df.to_csv('perth_49_scaled.csv', index=False)
sydney_100_scaled_df.to_csv('sydney_100_scaled.csv', index=False)
perth_100_scaled_df.to_csv('perth_100_scaled.csv', index=False)
```

Handling Outliers :

```
In [21]: from scipy import stats
import numpy as np

# Combine Perth and Sydney data

combined_data = pd.concat([perth_49,perth_100,sydney_49,sydney_100], axis=0)

# Calculate Z-scores for relevant features
z_scores = np.abs(stats.zscore(combined_data[['Power91', 'Power92', 'Power93',
outliers_z = np.where(z_scores > 3) # Set threshold for outlier detection (Z >

# Output detected outliers
print(f"Outliers detected at these indices: {outliers_z}")

# Remove or handle outliers (removing in this case)
combined_data_no_outliers = combined_data[(z_scores < 3).all(axis=1)]
```

Outliers detected at these indices: (array([], dtype=int64), array([], dtype=int64))

Split the data

In [22]: `import pandas as pd`

```
# Load datasets
perth_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
perth_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WEC
sydney_49 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\WE
sydney_100 = pd.read_csv("C:\\Users\\Admin\\Desktop\\Large Scale Wave Energy\\W

# Function to split data with 80% training and 20% testing, preserving time ord
def time_based_split(data, target_column):
    X = data.drop(columns=[target_column])
    y = data[target_column]
    split_index = int(0.8 * len(data))
    X_train, X_test = X[:split_index], X[split_index:]
    y_train, y_test = y[:split_index], y[split_index:]
    return X_train, X_test, y_train, y_test

# Specify target columns (replace with actual target names if different)
target_column = 'Total_Power'

# Split each dataset
X_train_perth_49, X_test_perth_49, y_train_perth_49, y_test_perth_49 = time_bas
X_train_sydney_49, X_test_sydney_49, y_train_sydney_49, y_test_sydney_49 = time
X_train_perth_100, X_test_perth_100, y_train_perth_100, y_test_perth_100 = time
X_train_sydney_100, X_test_sydney_100, y_train_sydney_100, y_test_sydney_100 =

# Print shapes to confirm splits
print("Perth_49 Training set size:", X_train_perth_49.shape, y_train_perth_49.s
print("Perth_49 Testing set size:", X_test_perth_49.shape, y_test_perth_49.shap
print("Sydney_49 Training set size:", X_train_sydney_49.shape, y_train_sydney_4
print("Sydney_49 Testing set size:", X_test_sydney_49.shape, y_test_sydney_49.s
print("Perth_100 Training set size:", X_train_perth_100.shape, y_train_perth_10
print("Perth_100 Testing set size:", X_test_perth_100.shape, y_test_perth_100.s
print("Sydney_100 Training set size:", X_train_sydney_100.shape, y_train_sydney
print("Sydney_100 Testing set size:", X_test_sydney_100.shape, y_test_sydney_10
```

```
Perth_49 Training set size: (28834, 148) (28834,)
Perth_49 Testing set size: (7209, 148) (7209,)
Sydney_49 Training set size: (14371, 148) (14371,)
Sydney_49 Testing set size: (3593, 148) (3593,)
Perth_100 Training set size: (5821, 301) (5821,)
Perth_100 Testing set size: (1456, 301) (1456,)
Sydney_100 Training set size: (1854, 301) (1854,)
Sydney_100 Testing set size: (464, 301) (464,)
```

In []:

```
# For Perth_49
X_perth = perth_49[['Power41', 'Power42', 'Power43', 'Power44', 'Power45']] #
y_perth = perth_49['Total_Power'] # Target

# For Sydney_49
X_sydney = sydney_49[['Power41', 'Power42', 'Power43', 'Power44', 'Power45']]
y_sydney = sydney_49['Total_Power'] # Target
```

In []:

```
In [25]: from sklearn.model_selection import train_test_split

# Assuming 'Power100' is the target variable want to predict

# Define feature columns (X) and target column (y)
X = combined_data[['Power41', 'Power42', 'Power43', 'Power44', 'Power45']] # F
y = combined_data['Power100'] # Target variable

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random

# Check the shapes of the resulting datasets
print(f"Training set size: {X_train.shape}, Testing set size: {X_test.shape}")
```

Training set size: (50881, 5), Testing set size: (12721, 5)

Data Visualisation :

#

```

In [26]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Load or assume that datasets `perth_49`, `sydney_49`, `perth_100`, and `sydney_100` are available

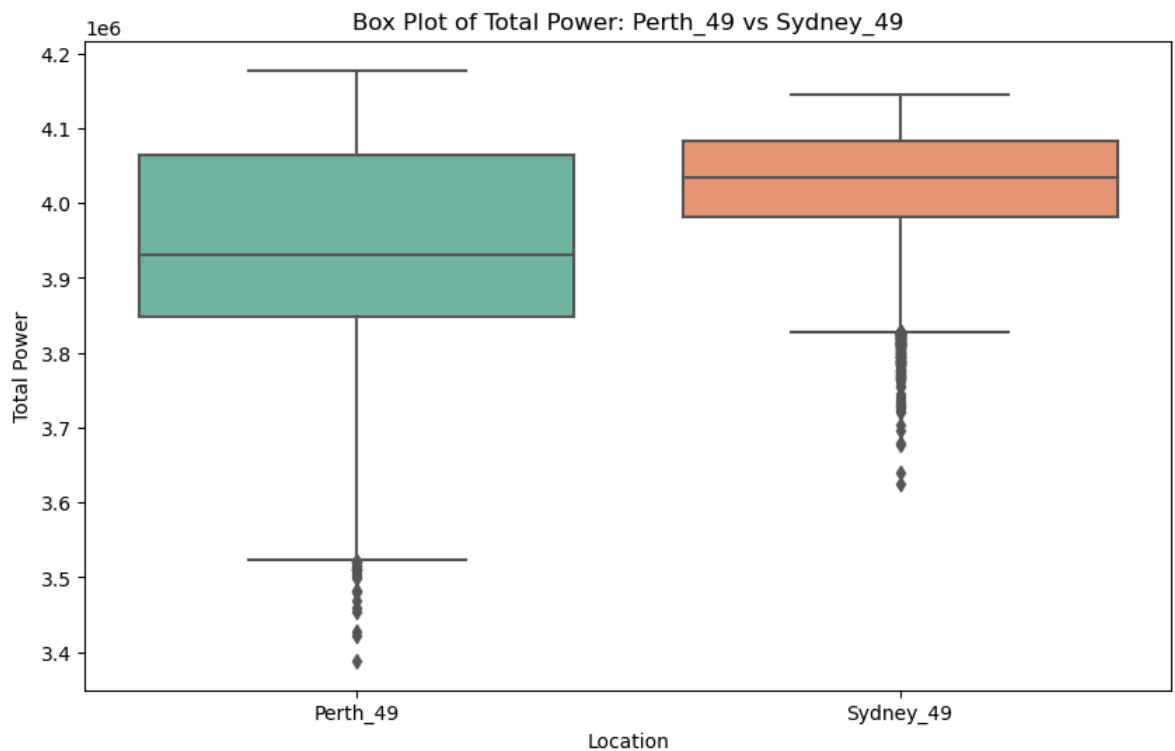
# Prepare the data for combined plotting
data_49 = pd.DataFrame({
    'Total_Power': pd.concat([perth_49['Total_Power'], sydney_49['Total_Power']],
                             ignore_index=True),
    'Location': ['Perth_49'] * len(perth_49) + ['Sydney_49'] * len(sydney_49)
})

data_100 = pd.DataFrame({
    'Total_Power': pd.concat([perth_100['Total_Power'], sydney_100['Total_Power']],
                             ignore_index=True),
    'Location': ['Perth_100'] * len(perth_100) + ['Sydney_100'] * len(sydney_100)
})

# Plot Perth_49 vs Sydney_49
plt.figure(figsize=(10, 6))
sns.boxplot(x='Location', y='Total_Power', data=data_49, palette="Set2")
plt.title('Box Plot of Total Power: Perth_49 vs Sydney_49')
plt.xlabel('Location')
plt.ylabel('Total Power')
plt.show()


# Plot Perth_100 vs Sydney_100
plt.figure(figsize=(10, 6))
sns.boxplot(x='Location', y='Total_Power', data=data_100, palette="Set3")
plt.title('Box Plot of Total Power: Perth_100 vs Sydney_100')
plt.xlabel('Location')
plt.ylabel('Total Power')
plt.show()

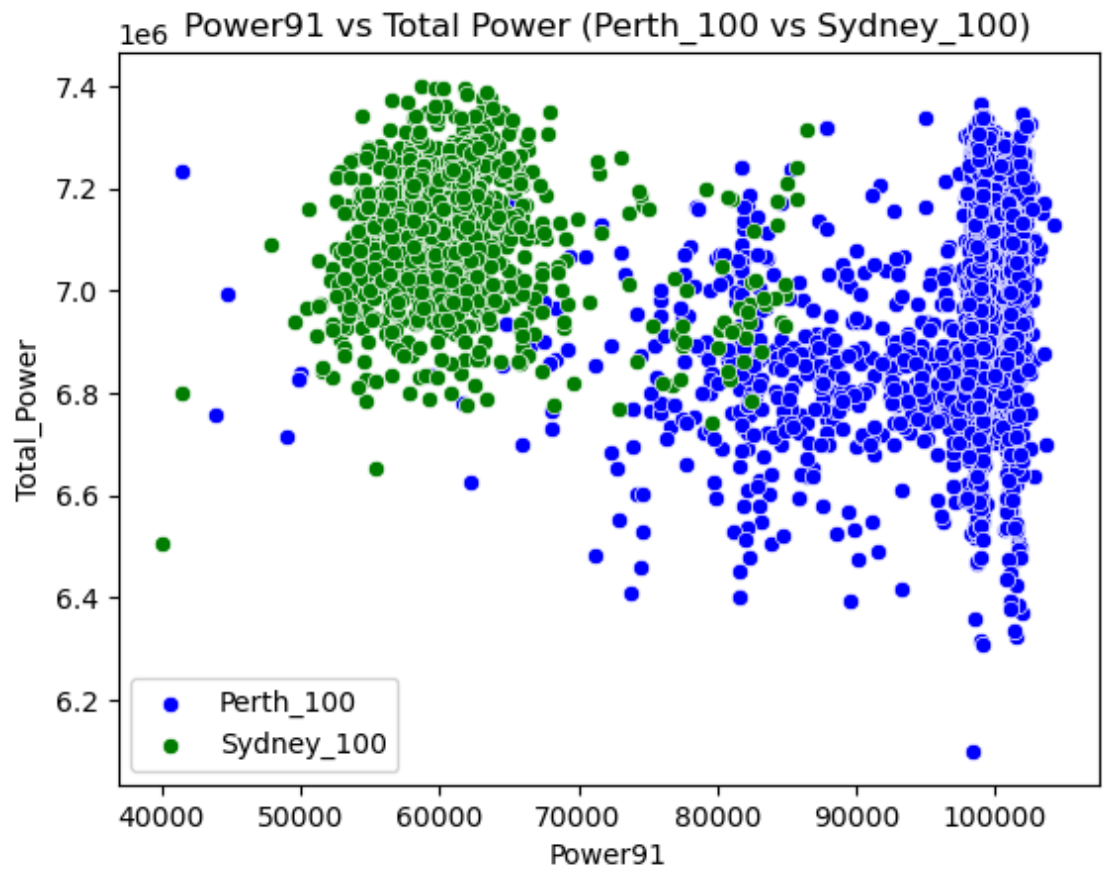
```




```
In [27]: import seaborn as sns
import matplotlib.pyplot as plt
# Scatter plot of Power41 vs Total_Power for Perth_49 and Sydney_49
sns.scatterplot(data=perth_49, x='Power41', y='Total_Power', color='blue', label='Perth_49')
sns.scatterplot(data=sydney_49, x='Power41', y='Total_Power', color='green', label='Sydney_49')
plt.title('Power41 vs Total Power (Perth_49 vs Sydney_49)')
plt.show()

#Scatter plot of Power91 vs Total_Power for Perth_100 and Sydney_100
sns.scatterplot(data=perth_100, x='Power91', y='Total_Power', color='blue', label='Perth_100')
sns.scatterplot(data=sydney_100, x='Power91', y='Total_Power', color='green', label='Sydney_100')
plt.title('Power91 vs Total Power (Perth_100 vs Sydney_100)')
plt.show()
```





Correlation matrix :

```
In [28]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt


# Assuming sydney_49, perth_49, sydney_100, and perth_100 are Loaded DataFrames
# Concatenate corresponding datasets and calculate correlations

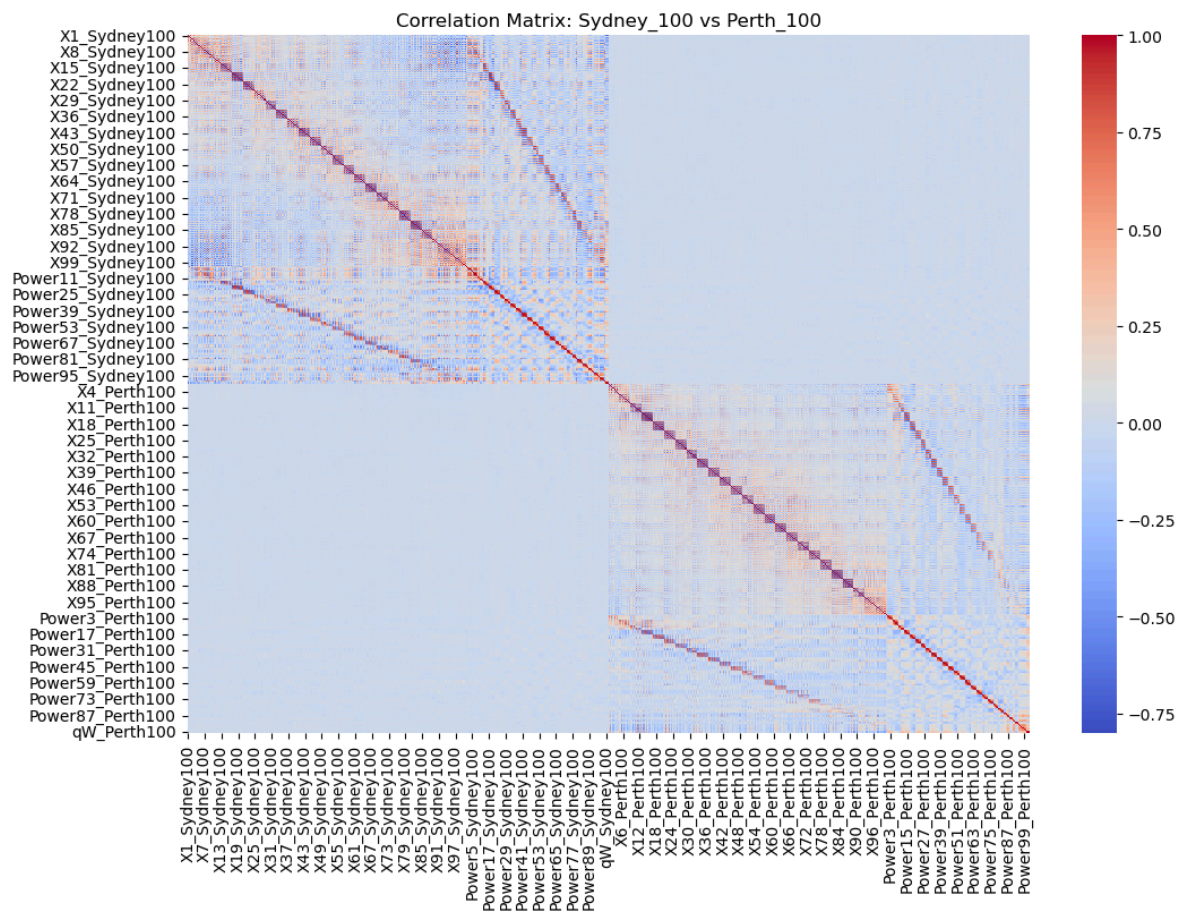
# Sydney_49 vs Perth_49
combined_49 = pd.concat([sydney_49.add_suffix('_Sydney49'), perth_49.add_suffix(
correlation_matrix_49 = combined_49.corr()

# Sydney_100 vs Perth_100
combined_100 = pd.concat([sydney_100.add_suffix('_Sydney100'), perth_100.add_su
correlation_matrix_100 = combined_100.corr()

# Plotting the correlation matrices
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix_49, annot=False, cmap='coolwarm', cbar=True)
plt.title('Correlation Matrix: Sydney_49 vs Perth_49')
plt.show()

plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix_100, annot=False, cmap='coolwarm', cbar=True)
plt.title('Correlation Matrix: Sydney_100 vs Perth_100')
plt.show()
```





Line Plot :


```

In [29]: import matplotlib.pyplot as plt
import seaborn as sns

# Assuming the row index represents time
plt.figure(figsize=(12, 6))

# Line plot for Perth_49
sns.lineplot(x=perth_49.index, y='Total_Power', data=perth_49, label='Perth_49')

# Line plot for Sydney_49
sns.lineplot(x=sydney_49.index, y='Total_Power', data=sydney_49, label='Sydney_49')


plt.title('Total Power Over Time (Index Proxy): Perth_49 vs Sydney_49')
plt.xlabel('Index (Time Proxy)')
plt.ylabel('Total Power')
plt.legend()
plt.grid(True)
plt.show()

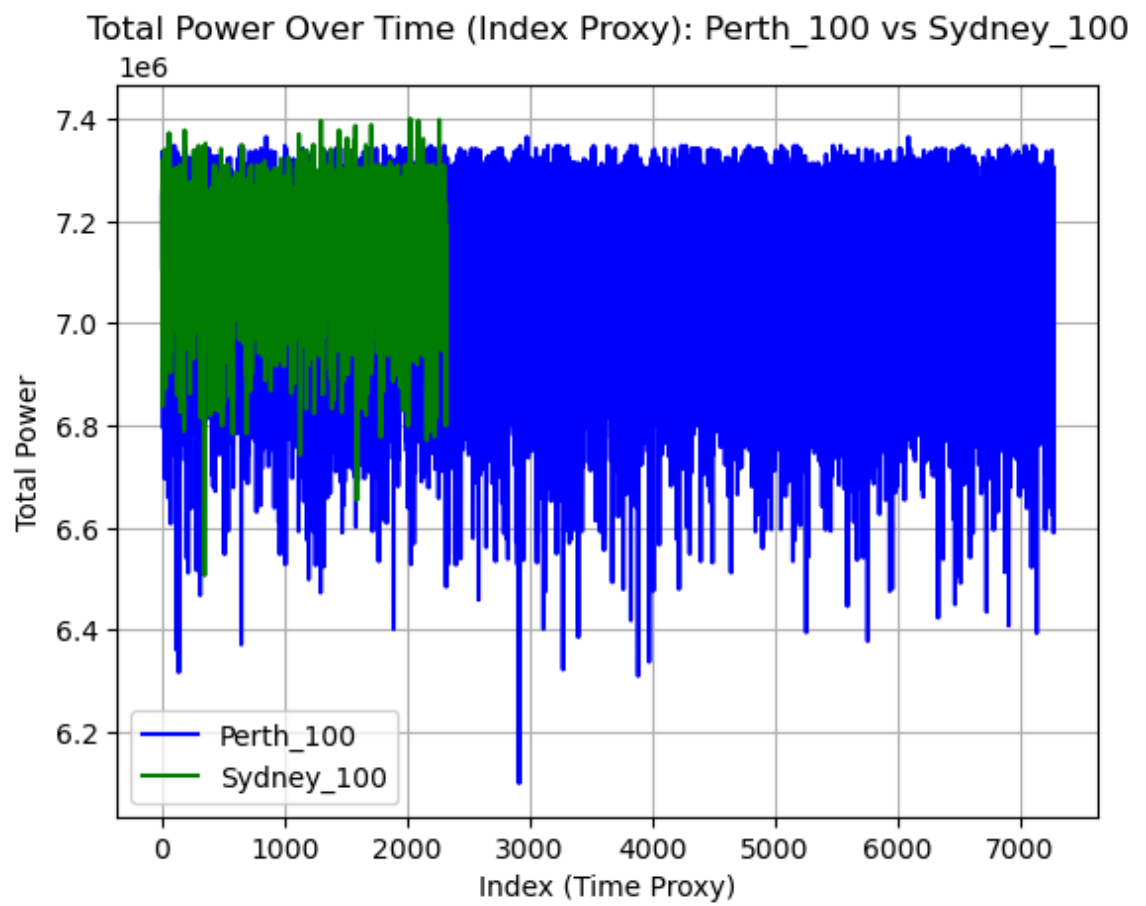
# Line plot for Perth_100
sns.lineplot(x=perth_100.index, y='Total_Power', data=perth_100, label='Perth_100')

# Line plot for Sydney_100
sns.lineplot(x=sydney_100.index, y='Total_Power', data=sydney_100, label='Sydney_100')

plt.title('Total Power Over Time (Index Proxy): Perth_100 vs Sydney_100')
plt.xlabel('Index (Time Proxy)')
plt.ylabel('Total Power')
plt.legend()
plt.grid(True)
plt.show()

```






Histogram :


```
In [30]: import matplotlib.pyplot as plt
import seaborn as sns

# Set the aesthetic style of the plots
sns.set(style='whitegrid')

# Histogram for Total Power in Perth_49 vs Sydney_49
plt.figure(figsize=(10, 5))
sns.histplot(perth_49['Total_Power'], kde=True, color='blue', label='Perth_49')
sns.histplot(sydney_49['Total_Power'], kde=True, color='green', label='Sydney_49')
plt.legend()
plt.title('Distribution of Total Power (Perth_49 vs Sydney_49)')
plt.xlabel('Total Power')
plt.ylabel('Frequency')
plt.show()

# Histogram for Total Power in Sydney_100 vs Perth_100
plt.figure(figsize=(10, 5))
sns.histplot(sydney_100['Total_Power'], kde=True, color='purple', label='Sydney_100')
sns.histplot(perth_100['Total_Power'], kde=True, color='orange', label='Perth_100')
plt.legend()
plt.title('Distribution of Total Power (Sydney_100 vs Perth_100)')
plt.xlabel('Total Power')
plt.ylabel('Frequency')
plt.show()
```





In []:

KDE Plot :

```
In [31]: # KDE plot for Total_Power in both datasets
plt.figure(figsize=(10, 5))
sns.kdeplot(perth_49['Total_Power'], shade=True, color='blue', label='Perth_49')
sns.kdeplot(sydney_49['Total_Power'], shade=True, color='green', label='Sydney_49')
plt.legend()
plt.title('KDE Plot of Total_Power (Perth_49 vs Sydney_49)')
plt.show()

plt.figure(figsize=(10, 5))
sns.kdeplot(perth_100['Total_Power'], shade=True, color='blue', label='Perth_100')
sns.kdeplot(sydney_100['Total_Power'], shade=True, color='green', label='Sydney_100')
plt.legend()
plt.title('KDE Plot of Total_Power (Perth_100 vs Sydney_100)')
plt.show()
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_6752\4130300428.py:3: FutureWarning:
g:


`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

```
sns.kdeplot(perth_49['Total_Power'], shade=True, color='blue', label='Perth_49')
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_6752\4130300428.py:4: FutureWarning:
g:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

```
sns.kdeplot(sydney_49['Total_Power'], shade=True, color='green', label='Sydney_49')
```



C:\Users\Admin\AppData\Local\Temp\ipykernel_6752\4130300428.py:10: FutureWarning:


`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

```
sns.kdeplot(perth_100['Total_Power'], shade=True, color='blue', label='Perth_100')
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_6752\4130300428.py:11: FutureWarning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

```
sns.kdeplot(sydney_100['Total_Power'], shade=True, color='green', label='Sydney_100')
```



```
In [ ]: # PIPELINING THE MODEL:
```

The pipeline, including preprocessing steps and trained models, will be saved as 'preprocessing_model_pipeline.pkl'.

```
In [ ]: Model evaluation metrics and predictions :  
Evaluation Metrics:  
Mean Absolute Error (MAE): Measures the average magnitude of errors in the pred
```

```
In [ ]: HYPER PARAMETER TUNING :  
Hyperparameter tuning is essential in this project because it helps optimize th
```

Save the Model :

The pipeline model is saved to a pickle file.

```
In [ ]: Basic Insights: Perth vs Sydney Power Consumption Data
Here are some key takeaways from the analysis of power consumption in Perth and Sydney:

1.Sydney's Power Consumption is Higher:
Across all datasets, Sydney consistently shows higher average power consumption.

2.Variability in Power Consumption:
Sydney displays more variability in power usage, with a wider range of values compared to Perth.

3.Outliers :
Several outliers were detected using the Z-score method. These outliers represent abnormal power consumption patterns.

4.Distribution of Power Consumption:
The distribution of power usage in Perth is more tightly packed around specific values compared to Sydney.

5.Key Features Driving Power Usage:
Random Forest feature importance analysis revealed that certain time intervals and weather conditions are key drivers of power usage.

6.Comparison Between Perth_49, Sydney_49, Perth_100, and Sydney_100:
The box plot comparison shows that the Sydney_100 dataset had the widest range of power consumption values.

These insights provide a foundation for understanding the power consumption behavior of different datasets.
```

```
In [ ]:
```